



Optimistic Concurrency Control in a Distributed NameNode Architecture for Hadoop Distributed File System

Qi Qi

Thesis to obtain the Master of Science Degree in
Information Systems and Computer Engineering

Supervisor: Doctor Luís Manuel Antunes Veiga

Examination Committee

Chairperson:	Doctor Luís Eduardo Teixeira Rodrigues
Supervisor:	Doctor Luís Manuel Antunes Veiga
Member of the Committee:	Doctor Nuno Preguiça

September 2014

Acknowledgments

The work presented is delivered as final thesis report at Instituto Superior Técnico - IST (Lisbon, Portugal). It is in partial fulfillment of the European Master in Distributed Computing - EMDC program 2012-2014. Royal Institute of Technology - KTH (Stockholm, Sweden) is the coordinator for this Erasmus Mundus master program. The study track has been composed of a first two semesters at IST, 3rd semester at KTH, and for this work and 4th semester, a degree project in Computer Systems Laboratory at Swedish Institute of Computer Science - SICS (Stockholm, Sweden).

Special thanks to my advisor Dr. Jim Dowling for his support throughout the project. With more than ten years' professional industry experience, Jim is always patient to help. He's the cool guy who gives answers faster than Google and StackOverflow.

Thanks to Salman Niazi and Mahmoud Ismail for all the practical help. Without them I might have to spend quite a long time studying the code base of the precedent work.

I'm also grateful to my supervisor Prof. Luís Antunes Veiga for his continuous support and encouragement. When I was in IST, I liked staying in the classroom after his class and chatted with him for a while. Veiga was like a big brother there taking care of us.

I would like to thank the good friends I met in Portugal and Sweden, who leveled me up during these two years. Without you guys, this journey wouldn't have been such a legendary in my life.

I am truly thankful to my family for nursing me with all their affections and love.

Last, special appreciation to this young man, Qi Qi, who always has the guts to go on any adventure in his life.

September 8, 2014, Stockholm

Qi Qi

Dedication

*To my father, a man of integrity, who
supports all my adventurous decisions so
that I can live outside of the box.*

Resumo

[To be added] Portuguese Abstract

Abstract

The *Hadoop Distributed File System* (HDFS) is the storage layer for Apache Hadoop ecosystem, persisting large data sets across multiple machines. However, the overall storage capacity is limited since the metadata is stored in-memory on a single server, called the *NameNode*. The heap size of the *NameNode* restricts the number of data files and addressable blocks persisted in the file system.

The *Hadoop Open Platform-as-a-service* (Hop) is an open platform-as-a-Service (PaaS) support of the Hadoop ecosystem on existing cloud platforms including Amazon Web Service and OpenStack. The storage layer of Hop, called the Hop-HDFS, is a highly available implementation of HDFS, based on storing the metadata in a distributed, in-memory, replicated database, called the *MySQL Cluster*. It aims to overcome the *NameNode*'s limitation while maintaining the strong consistency semantics of HDFS so that applications written for HDFS can run on Hop-HDFS without modifications.

Precedent thesis works have contributed for a transaction model for Hop-HDFS. From system-level coarse grained locking to row-level fine grained locking, the strong consistency semantics have been ensured in Hop-HDFS, but the overall performance is restricted compared to the original HDFS.

In this thesis, we first analyze the limitation of HDFS *NameNode* implementation and provide an overview of Hop-HDFS illustrating how we overcome those problems. Then we give a systematic assessment on precedent works for Hop-HDFS comparing to HDFS, and also analyze the restriction when using pessimistic locking mechanisms to ensure the strong consistency semantics. Finally, based on the investigation of current shortcomings, we demonstrate how to improve the performance by designing a new model based on optimistic concurrency control with snapshot isolation. As a proof of concept, the evaluation shows the significant improvement of this new model. The correctness of our implementation has been validated by 300+ Apache HDFS unit tests passing.

Palavras Chave

Keywords

Palavras Chave [To be corrected by native Portuguese speaker]

HDFS

MySQL Cluster

Controle de Concorrência

Snapshot Isolation

Transação

Vazão

Keywords

HDFS

MySQL Cluster

Concurrency Control

Snapshot Isolation

Transaction

Throughput

Index

I	Introduction and Background	1
1	Introduction	3
1.1	Motivation	3
1.1.1	The De Facto Industrial Standard in Big Data Era	3
1.1.2	Limits to growth in HDFS	3
1.1.3	Hop-HDFS and Its Limitation	4
1.2	Problem Statement	5
1.3	Contribution	5
1.4	Document Structure	5
2	Background and Related Work	7
2.1	A	7
2.2	B	7
2.3	C	7
2.4	D	7
II	Assessment in Hop-HDFS	9
3	Limitation on Pessimistic Locking Mechanism	11
3.1	A	11

3.2	B	11
3.2.1	B1	11
3.2.2	B2	11
3.3	C	11
3.4	D	11
4	Systematic Assessment of Hop-HDFS Performance	13
4.1	A	13
4.2	B	13
4.2.1	B1	13
4.2.2	B2	13
4.3	C	13
4.4	D	13
III	Solution	15
5	Design	17
5.1	A	17
5.2	B	17
5.2.1	B1	17
5.2.2	B2	17
5.3	C	17
5.4	D	17
6	Implementation	19
6.1	A	19

6.2	B	19
6.2.1	B1	19
6.2.2	B2	19
6.3	C	19
6.4	D	19
IV	Evaluation and Conclusion	21
7	Evaluation	23
7.1	A	23
7.2	B	23
7.2.1	B1	23
7.2.2	B2	23
7.3	C	23
7.4	D	23
8	Conclusion	25
8.1	A	25
8.2	B	25
8.2.1	B1	25
8.2.2	B2	25
8.3	C	25
8.4	D	25
V	Appendices	29
A	Apache HDFS Unit Tests Passing List	31

List of Figures

List of Tables

1.1	Memory Requirement for Related Storage Capacity in HDFS	4
-----	---	---



Introduction and Background

1

Introduction

1.1 Motivation

1.1.1 The De Facto Industrial Standard in Big Data Era

The *Apache Hadoop* ([Apache](#)) ecosystem has become the de facto industrial standard to store, process and analyze large data sets in the big data era ([Cloudera](#)). It is widely used as a computational platform for a variety of areas including search engines, data warehousing, behavioral analysis, natural language processing, genomic analysis, image processing, etc ([Shvachko 2011](#)).

The *Hadoop Distributed File System* (HDFS) is the storage layer for Apache Hadoop, which enables petabytes of data to be persisted on clusters of commodity hardware at relatively low cost ([Borthakur 2008](#)). Inspired by the *Google File System* (GFS) ([Ghemawat et al. 2003](#)), the namespace, *metadata*, is decoupled from data and stored in-memory on a single server, called the *NameNode*. The file datasets are stored as sequences of blocks and replicated across potentially thousands of machines for fault tolerance.

1.1.2 Limits to growth in HDFS

Built upon the single namespace server, *the NameNode*, architecture, one well-known limitation of HDFS is the limitation to growth ([Shvachko 2010](#)). Since the metadata is kept in-memory for fast operation in NameNode, the number of file objects in the filesystem is limited by the amount of memory of the NameNode.

Approximately, the size of the metadata for a single file object having two blocks (replicated three times by default) is 600 bytes. As a rule of thumb, for one petabyte physical storage, it requires one gigabyte metadata in memory ([Shvachko 2010](#)). Table 1.1 gives an estimation of the memory requirement and its related physical storage capacity for different number of files.

Number of Files	Memory Requirement	Physical Storage
1 million	0.6 GB	0.6 PB
100 million	60 GB	60 PB
1 billion	600 GB	600 PB
2 billion	1200 GB	1200 PB

Table 1.1: Memory Requirement for Related Storage Capacity in HDFS

As HDFS runs in the *Java Virtual Machine* (JVM), due to interactive workloads, heap sizes larger than 60 GB is not considered practical (Shvachko 2010). Therefore, 100 million files will be the maximum storage capacity of HDFS.

1.1.3 Hop-HDFS and Its Limitation

The *Hadoop Open Platform-as-a-service* (Hop) (Dowling 2013) is an open platform-as-a-Service (PaaS) support of the Hadoop ecosystem on existing cloud platforms including Amazon Web Service and OpenStack. The storage layer of Hop, called the Hop-HDFS, is a highly available implementation of HDFS, based on storing the metadata in a distributed, in-memory, replicated database, called the *MySQL Cluster*. It aims to overcome the NameNode's limitation while maintaining the strong consistency semantics of HDFS so that applications written for HDFS can run on Hop-HDFS without modifications.

Precedent thesis works have contributed for a transaction model (Wasif 2012) (Peiro Sajjad & Hakimzadeh Harirbaf 2013) as well as a high availability multi-NameNode architecture (D'Souza 2013) for Hop-HDFS. It can store up to 4.1 billion files with 3TB MySQL Cluster support for metadata (Hakimzadeh et al. 2014).

However, in HDFS, the correctness and consistency of the namespace is ensured by atomic metadata mutation (Shvachko et al. 2010). In order to maintain the same level of strong consistency semantics, system-level coarse grained locking and row-level fine grained locking are adopted in precedent projects of Hop-HDFS, but the overall performance is heavily restricted compared to the original HDFS. Therefore, investigation for better concurrency control to improve the performance of Hop-HDFS is the main motivation.

1.2 Problem Statement

In HDFS, the NameNode's operations are categorized into *read* or *write* operations. To protect the metadata among parallel running threads, a global read/write lock is used to maintain the atomicity of the namespace. We call it *system-level lock*. Concurrent threads to access shared object for read operations are allowed, but it restricts a single thread to access object for write operations. This *single-writer-multiple-readers* concurrency model will not reduce the throughput much since the metadata is kept optimized data structures in-memory (Hakimzadeh et al. 2014).

The first version of Hop-HDFS, called the KTHFS (Wasif 2012), adopts the system-level locking mechanism to serialize transactions. All concurrent readers get the same view of the mutated data reflected by completed writes. We call it *Strong Consistency*.

1.3 Contribution

CCC

1.4 Document Structure

[To be Added]

Background and Related Work

2.1 *A*

AAA

2.2 *B*

BBB

2.3 *C*

CCC

2.4 *D*

DDD

II Assessment in Hop-HDFS

3

Limitation on Pessimistic Locking Mechanism

3.1 *A*

AAA

3.2 *B*

BBB

3.2.1 **B1**

BBB1

3.2.2 **B2**

BBB2

3.3 *C*

CCC

3.4 *D*

DDD

4 Systematic Assessment of Hop-HDFS Performance

Neque porro quisquam est qui dolorem ipsum quia dolor sit amet, consectetur, adipisci velit...

– Cerico

4.1 *A*

AAA

4.2 *B*

BBB

4.2.1 **B1**

BBB1

4.2.2 **B2**

BBB2

4.3 *C*

CCC

4.4 *D*

DDD

III

Solution

5

Design

5.1 *A*

AAA

5.2 *B*

BBB

5.2.1 **B1**

BBB1

5.2.2 **B2**

BBB2

5.3 *C*

CCC

5.4 *D*

DDD

6

Implementation

6.1 *A*

AAA

6.2 *B*

BBB

6.2.1 **B1**

BBB1

6.2.2 **B2**

BBB2

6.3 *C*

CCC

6.4 *D*

DDD

IV

Evaluation and Conclusion

7

Evaluation

7.1 *A*

AAA

7.2 *B*

BBB

7.2.1 **B1**

BBB1

7.2.2 **B2**

BBB2

7.3 *C*

CCC

7.4 *D*

DDD

8

Conclusion

8.1 *A*

AAA

8.2 *B*

BBB

8.2.1 **B1**

BBB1

8.2.2 **B2**

BBB2

8.3 *C*

CCC

8.4 *D*

DDD

Bibliography

- Apache. What is apache hadoop? <http://hadoop.apache.org>.
- Borthakur, D. (2008). Hdfs architecture guide. *HADOOP APACHE PROJECT* [http://hadoop. apache. org/common/docs/current/hdfs design. pdf](http://hadoop.apache.org/common/docs/current/hdfs design. pdf).
- Cloudera. Hadoop and big data. <http://www.cloudera.com/content/cloudera/en/about/hadoop-and-big-data.html>.
- Dowling, J. (2013). Hop: Hadoop open platform-as-a-service.
- D’Souza, J. C. (2013). Kthfs—a highly available andscalable file system.
- Ghemawat, S., H. Gobioff, & S.-T. Leung (2003). The google file system. In *ACM SIGOPS Operating Systems Review*, Volume 37, pp. 29–43. ACM.
- Hakimzadeh, K., H. P. Sajjad, & J. Dowling (2014). Scaling hdfs with a strongly consistent relational model for metadata. In *Distributed Applications and Interoperable Systems*, pp. 38–51. Springer.
- Peiro Sajjad, H. & M. Hakimzadeh Harirbaf (2013). Maintaining strong consistency semantics in a horizontally scalable and highly available implementation of hdfs.
- Shvachko, K., H. Kuang, S. Radia, & R. Chansler (2010). The hadoop distributed file system. In *Mass Storage Systems and Technologies (MSST), 2010 IEEE 26th Symposium on*, pp. 1–10. IEEE.
- Shvachko, K. V. (2010). Hdfs scalability: The limits to growth. *login* 35(2), 6–16.
- Shvachko, K. V. (2011). Apache hadoop: The scalability update. *login: The Magazine of USENIX* 36, 7–13.
- Wasif, M. (2012). A distributed namespace for a distributed file system.



Appendices



Apache HDFS Unit Tests Passing List

